

# NeSy-Core — Production File Structure

```
nesy-core/
|
├── README.md
├── LICENSE
├── pyproject.toml          # Build system config (PEP 517)
├── setup.cfg               # Package metadata
├── Makefile                # Dev shortcuts: make test, make lint, make b
├── .env.example             # Environment variables template
├── .gitignore
├── .pre-commit-config.yaml # Code quality hooks
|
└── nesy/                   # ← Main Python package
    ├── __init__.py
    └── version.py           # Single source of truth for version
|
    ├── core/                # Core framework – no ML dependencies
    │   ├── __init__.py
    │   ├── types.py
    │   ├── exceptions.py
    │   ├── config.py
    │   └── registry.py       # Component registry for plugins
    |
    ├── symbolic/             # Symbolic reasoning engine
    │   ├── __init__.py
    │   ├── engine.py
    │   ├── logic.py
    │   ├── rules.py
    │   ├── solver.py
    │   ├── betti.py
    │   └── ontology/
    │       ├── __init__.py
    │       ├── loader.py      # Load OWL, RDF, custom ontologies
    │       └── adapters/
    │           ├── owl.py
    │           └── rdf.py
    |
    └── neural/                # Neural backbone
        ├── __init__.py
        ├── base.py
        └── grounding.py      # Symbol grounding: embeddings → predicates
```



```
|   └── configs/                      # Domain-specific configs
|       ├── default.yaml
|       ├── medical.yaml
|       ├── legal.yaml
|       └── edge.yaml
|
|   └── tests/                         # Full test suite
|       ├── conftest.py                # Shared fixtures
|       └── unit/
|           ├── test_symbolic_engine.py
|           ├── test_logic.py
|           ├── test_grounding.py
|           ├── test_metacognition.py
|           ├── test_confidence.py
|           └── test_continual.py
|       └── integration/
|           ├── test_pipeline.py        # End-to-end pipeline tests
|           ├── test_huggingface.py
|           └── test_server.py
|   └── benchmarks/
|       ├── test_inference_speed.py
|       ├── test_memory_usage.py
|       └── test_npu_latency.py
|
|   └── examples/                      # Working code examples
|       ├── basic_reasoning.py
|       ├── medical_diagnosis.py
|       ├── continual_learning.py
|       ├── edge_deployment.py
|       └── notebooks/
|           ├── 01_quickstart.ipynb
|           ├── 02_symbolic_rules.ipynb
|           └── 03_metacognition.ipynb
|
|   └── docs/                           # Documentation source
|       ├── index.md
|       ├── quickstart.md
|       ├── architecture.md
|       └── api_reference/
|           ├── nesy_model.md
|           ├── symbolic_engine.md
|           └── metacognition.md
|   └── guides/
|       ├── adding_rules.md
|       ├── edge_deployment.md
|       └── custom_backbone.md
```

```

|   └── scripts/                      # Dev and ops scripts
|       ├── build_concept_graph.py     # Build domain concept graph from corpus
|       ├── evaluate_benchmarks.py    # Run full benchmark suite
|       ├── export_model.py          # Export trained model
|       └── generate_docs.py

|
|   └── docker/
|       ├── Dockerfile                # Production image
|       ├── Dockerfile.dev           # Development image
|       └── docker-compose.yml       # Full stack: server + dependencies

|
└── .github/
    ├── workflows/
    |   ├── ci.yml                  # Run tests on every PR
    |   ├── release.yml            # Publish to PyPI on tag
    |   └── benchmark.yml         # Performance regression checks
    └── ISSUE_TEMPLATE/
        ├── bug_report.md
        └── feature_request.md

```

---

## Key Design Decisions

**Why `nesy/core/types.py` exists separately:** Every module imports from `core/types.py`. No circular imports. Clean dependency graph.

**Why `symbolic/` and `neural/` are siblings, not parent/child:** Either can work without the other. A pure symbolic pipeline is valid. A pure neural pipeline is valid. The `api/bridge.py` connects them — but does not own either.

**Why `metacognition/` is its own top-level module:** It is the primary novelty. It should be importable independently: `from nesy.metacognition import MetaCognitionMonitor`. Teams can adopt it without the full framework.

**Why `continual/symbolic_anchor.py` is separate from the neural weights:** Symbolic facts are immutable. Neural weights drift. They must never be in the same storage layer.

**Why `deployment/` has its own `server/`:** The inference server is production code, not a script. It needs its own routes, middleware, and auth — treated as a microservice.

---

## Install Targets

```
pip install nesy-core                      # Core only
pip install nesy-core[torch]                  # + PyTorch backbone
pip install nesy-core[server]                # + FastAPI server
pip install nesy-core[edge]                 # + ONNX + TFLite export
pip install nesy-core[all]                   # Everything
```

---

## First File to Write

Start here: nesy/core/types.py

```
from dataclasses import dataclass, field
from typing import List, Dict, Optional
from enum import Enum

class ConfidenceType(Enum):
    FACTUAL = "factual"
    REASONING = "reasoning"
    KNOWLEDGE_BOUNDARY = "knowledge_boundary"

@dataclass
class SymbolicRule:
    id: str
    predicate: str
    constraints: List[str]
    weight: float = 1.0

@dataclass
class ReasoningTrace:
    steps: List[str]
    rules_activated: List[SymbolicRule]
    neural_confidence: float
    symbolic_confidence: float

@dataclass
class NSIOutput:
    answer: str
    confidence: Dict[ConfidenceType, float]
    reasoning_trace: ReasoningTrace
    flags: List[str] = field(default_factory=list)
    status: str = "ok"    # ok | flagged | uncertain
```

Everything else builds on top of these types.

